

GENERATIVE AGENTS IN AGENT-BASED MODELING

Overview, Validation and Emerging Challenges

Submitted by:

Aravind A Kamath (CEC23CS041)

Under the Guidance of:

Mrs. JAYASREE K

Assistant Professor

Department of Computer Science and Engineering
College of Engineering, Cherthala

Presentation Outline

- 1 Introduction and Theoretical Foundations
- 2 Evolution of Agents in ABM
- 3 Architecture and Mechanics of Generative Agents
- 4 Notable Implementations and Frameworks
- 5 Validation Methodologies for ABM and GABM
- 6 Comparative Analysis: Traditional ABM vs. GABM
- 7 Emerging Challenges in GABMs
- 8 Discussion and Proposed Hybrid Architecture
- 9 Conclusion and References

Background: Agent-Based Modeling (ABM)

Definition and Scope:

- Computational simulation of autonomous entities (agents) interacting locally.
- Primary objective: Observe emergent macroscopic patterns from micro-level decisions.

Traditional Agent Mechanics:

- Built on rigid mathematical equations, heuristics or 'if-then' rules.
- Assumes strict economic rationality (*homo economicus*).

Core Limitations:

- Fails to capture human bounded rationality, bias and emotional states.
- Cannot model natural language dialogue or open-ended reasoning.

The Concept of 'Generative' in Social Science

Historical Terminology (Epstein, 1999):

- Core principle: *"If you didn't grow it, you didn't explain it."*
- In classical ABM, 'generative' meant growing macro structures from local micro rules without central coordination.

Foundational Classical Models:

- **Schelling's Segregation (1971):** Mild individual neighbor preferences generate extreme macroscopic segregation.
- **Reynolds' Boids (1987):** Lifelike flocking emerges from three local rules (Separation, Alignment, Cohesion).

The Modern Terminological Shift:

- Classical generative models did **not** utilize generative artificial intelligence.
- Modern GABMs explicitly leverage Large Language Models (LLMs) to synthesize autonomous cognitive behaviors.

Traditional ABM vs. Generative ABM: Urban Traffic Flow

Table: Comparison of the same simulation scenario under ABM and GABM (Table I of base paper)

Dimension	Traditional ABM	Generative ABM (GABM)
Scenario	Traffic flow in a city with cars managed by signals and road closures.	Traffic flow incorporating human social norms, emotions and communication.
Objective	Optimizes physical throughput, predicts congestion bottlenecks, visually graphs flow.	Uncovers how social norms, personal risk, stress and individual decisions shape mobility patterns.
Outcomes & Generalization	Produces context-specific engineering outputs for urban planners; findings cannot generalize into broader social theories.	Generates theoretical insights into social behavior-infrastructure interplay; generalizable across domains.
Decision Basis	Purely rational shortest-path algorithms (Dijkstra/A*), ignoring subjective mental states.	Context-sensitive reasoning accounting for fear, comfort, scenic preference and social influence.

Evolution of Agents: Evolutionary Algorithms (EAs)

Darwinian Principles in ABM:

- Introduces adaptation through selection, mutation and crossover.
- Genetic Algorithms (GAs) and Evolution Strategies optimize agent parameters via fitness functions.

Cultural Algorithms and Belief Spaces:

- Maintains an explicit shared 'belief space' (population-level knowledge repository).
- Extracts collective experiences to guide individual agent evolution.

Key Limitations of EA Approaches:

- High computational cost from iterative multi-generation evaluations.
- Risk of premature convergence to suboptimal local equilibria.
- High sensitivity to hyperparameter tuning (mutation rates, selection pressure).

Evolution of Agents: Reinforcement Learning (RL)

Trial-and-Error Policy Optimization:

- Agents optimize actions to maximize cumulative reward signals.
- Applications: Pedestrian dynamics, evacuation routes and financial markets.

Key Challenges in Social Simulations:

- **Stochasticity & Suboptimality:** Early exploration causes erratic behaviors; easily trapped in local optima.
- **Brittle Generalization:** Policies struggle when deployed in novel, data-scarce social environments.
- **Lack of Language:** Optimizes numerical state tensors without support for semantic dialogue, negotiation or persuasion.

AI Integration in ABM Development

- Machine Learning integrates into ABM across two primary areas: **Structural Specification** and **Output Analysis**.




Figure: Key roles of Generative AI in ABM development (Source: Base Paper, Section III-A)

Defining Generative Agents (GAs)

- **Formal Definition:**

- A Generative Agent (GA) is an autonomous computational entity powered by a Large Language Model, capable of displaying human-like behavior, reasoning contextually and interacting via natural language.

- **The LLM Feedback Loop Concept:**

- Environmental feedback loop: The simulation environment and agents continuously exchange feedback regarding states and actions with the underlying LLM.
- **Mechanical Model:** Rule-based physical substrate defined by developers (coordinates, physics, obstacle collision, time progression).
- **Cognitive Layer:** Large Language Model interprets the sensory perception, queries its internal knowledge and memories and generates believable behavioral responses.

- **Key Breakthrough:**

- Behavioral rules emerge organically from pre-trained world knowledge rather than being rigidly pre-programmed by researchers.

GA Decision Architectures: Inferential vs. Conversational




Figure: LLM-based feedback loop architectures for Generative Agents (Source: Base Paper, Section IV)

- **Inferential GAs [Fig. 2(a)]:**

- LLM serves purely as a centralized decision engine.
- Inputs: Current Agent Status + Environmental Event → Output: Discrete Action.

- **Conversational GAs [Fig. 2(b)]:**

- Agents engage in open-ended natural language exchanges with peers.
- Dialogues directly alter internal beliefs, relationships and emotional statuses.

Internal State Representation: Parameters vs. Dynamic Status

Agent Parameters (Static Baseline):

- **Demographics:** Age, gender, income tier, occupation, education.
- **Psychological Profile:** Big Five personality traits (OCEAN).
- **Social Network:** Baseline relationships, family ties and roles.

Dynamic Agent Status (Runtime Representation):

- **State Representation:** Captures agent state during simulation runtime.
- **Decoupled Architecture:** Separates internal state from LLM decision logic.
- **Affective Tracking:** Records dynamic sentiment, emotional shifts and interaction history.

Core Cognitive Components: Memory Stream & Retrieval

- **Memory Stream (Park et al., 2023):**

- A comprehensive chronological database recording all sensory observations, natural language dialogues and internal thoughts.
- Overcomes LLM finite context window limits by externalizing long-term experience storage.

- **Contextual Multi-Factor Retrieval Function:**

- Whenever an agent decides or speaks, it queries its memory stream using a weighted linear combination of three normalized scores:

$$\text{Score} = \alpha \cdot \text{Recency} + \beta \cdot \text{Importance} + \gamma \cdot \text{Relevance}$$

- **Recency:** Exponential decay function prioritizing recent events over distant memories.
- **Importance:** Integer score assigned by LLM separating trivial routines (e.g., brushing teeth) from pivotal events (e.g., job offer).
- **Relevance:** Cosine similarity between embedding vector of current context and memory statement embedding.

Core Cognitive Components: Reflection & Dynamic Planning

● Reflection Hierarchy (Knowledge Synthesis):

- Raw sensory observations are insufficient for high-level social intelligence.
- The reflection module periodically triggers when accumulated importance scores cross a threshold.
- Queries the 100 most recent memories, synthesizes patterns and generates higher-level abstract thoughts.
- Forms an evolving directed tree: Base observations (leaf nodes) → Abstract beliefs and values (higher-level nodes).

● Planning and Reactive Execution Loop:

- **Top-Down Plan Decomposition:** Agents formulate realistic broad daily agendas, recursively decomposing them into hourly schedules, then minute-by-minute action sequences.
- **Dynamic Reactive Re-planning:** When an unexpected event or conversation occurs, the agent pauses its schedule, re-evaluates through the LLM and updates its plan adaptively.

Notable Implementations: The Smallville Sandbox

- **Experimental Sandbox (Park et al., 2023):**

- 2D sprite environment (resembling *The Sims*) populated by 25 autonomous generative agents.
- Agents initialized with simple text seed descriptions (name, occupation, basic relationships).

- **Emergent Macroscopic Phenomena:**

- **Information Diffusion:** When one agent announced a desire to host a Valentine's Day party, the news organically diffused across the town over two simulation days through peer-to-peer dialogues.
- **Autonomous Coordination:** Agents independently decorated the venue, invited friends and showed up at the party at the specified hour without hardcoded triggers.
- **Emergent Social Ties:** Agents formed new friendships and political opinions regarding town mayoral elections.

- **Ablation Findings:**

- Removing either Memory, Reflection or Planning severely degraded agent believability and long-term behavioral coherence.

General-Purpose Simulation Frameworks

- **Concordia (Google DeepMind / Vezhnevets et al., 2023):**
 - General-purpose simulation framework for language-mediated interactions in grounded digital, physical or social spaces.
 - Introduces a centralized '**Game Master**' agent: Grounds language actions into physical reality, enforcing physical constraints, resource budgets and institutional rules.
 - Bridges unconstrained LLM text generation with structured environmental state transitions.
- **SOTOPIA (Zhou et al., 2024):**
 - Open-ended interactive environment designed to evaluate social intelligence in language agents.
 - Simulates complex, goal-driven interpersonal scenarios involving cooperation, competition, deception and negotiation.
 - Features cross-agent profiles and provides quantitative benchmarks for multi-agent social dialogues.

- **CRSEC Architecture for Social Norm Emergence (2024):**

- Four-stage cognitive pipeline modeling how social norms naturally emerge and get enforced in GA societies:
 - ① **Creation & Representation:** Formalizes novel behavioral rules from observed social friction.
 - ② **Spreading:** Disseminates normative expectations through peer conversations.
 - ③ **Evaluation:** Agents assess conformity and compute social sanction penalties.
 - ④ **Compliance:** Internalizes rules into future planning routines.
- Successfully validated with 10 GAs in Smallville (e.g., establishing indoor smoking bans).

- **S3: Social-Network Simulation System (Gao et al., 2023):**

- Models digital information propagation, emotional contagion and echo chamber polarization across virtual online platforms.
- Empirically replicates real-world retweet cascades and opinion divergence.

The Fundamental Validation Dilemma

- **Why Validation is Essential:**

- Critical for establishing trust, credibility and ensuring simulation findings reflect real-world social processes rather than computational artifacts.

- **The Core Epistemological Hurdle:**

- Traditional ABM agents are transparent white-box models: Explicit rules can be mathematically inspected.
- Generative Agents are non-deterministic, probabilistic black boxes: Multiple runs with identical inputs produce divergent outputs.
- Internal reasoning paths cannot be uniformly traced back to closed-form equations.

- **Hierarchy of Evidence (inspired by Evidence-Based Medicine):**

- **Highest Rigor:** Ecological validity and controlled experimental trials.
- **Intermediate:** Observational field data matching.
- **Lowest Rigor:** Theoretical plausibility and expert consistency.

Traditional ABM Validation Methods

- **Foundational Validation Methods (Base Paper Section V-A):**

- **Data Analytics:** Statistical cleaning, feature correlation and input/output distribution verification.
- **Docking (Model Alignment):** Comparing simulation outputs against established baseline models under identical parameters.
- **Empirical Validation:** Matching simulated macroscopic outputs against observed historical datasets.
- **Sampling:** Sensitivity testing across parameter spaces to evaluate model robustness.

- **Advanced Validation Methods:**

- **Bootstrapping:** Resampling simulation runs to build empirical confidence intervals without assuming normal distributions.
- **Causal Analysis:** Evaluating intervention effects to distinguish genuine causality from spurious correlations.
- **Inverse Generative Social Science (IGSS):** Using genetic programming to backward-engineer agent rule structures that produce target macro-patterns.
- **Spatial Analytics:** Validating topological and geographic distribution

Face Validation: Qualitative and Human-Centric

- **Definition of Face Validation:**

- Assessment of whether model operational mechanics and outputs appear realistic, believable and plausible to human domain experts.

- **Microface Validation:**

- Evaluates individual agent behaviors at the micro-level.
- Verifies that personal choices, emotional transitions and natural language statements are logically coherent and believable over time.


- **Macroface Validation:**

- Examines aggregate, system-level simulation patterns (e.g., crowd evacuation flows, spatial segregation, traffic bottlenecks).
- Assesses whether collective emergent behaviors match intuitive real-world expectations.

- **Immersive Assessment:**

- Allows human experts to enter the virtual environment directly (via interactive UI or VR) to interview agents and evaluate responses in real time.

GABM Validation Workflows: The Integrated Architecture



- **Empirical Pathway:** Data analytics and docking for quantitative macro-matching.
- **Macroface Pathway:** Spatial and causal analysis for collective emergent patterns.
- **Microface Pathway:** Interviews, believability scoring and dialogue audits.

Validating Existing GAs: Empirical Implementations

- **Smallville Controlled Evaluation (Park et al., 2023):**

- Human evaluators interviewed agents across five core cognitive categories:
 - ① Self-Knowledge 2. Memory 3. Plans 4. Reactions 5. Reflections
- Evaluated 100 human judges across 5 ablation conditions to measure believability.

- **SOTOPIA-EVAL Benchmark (Zhou et al., 2024):**

- Quantitative evaluation framework measuring social intelligence across 7 dimensions:
 - **Goal Completion:** Objective task attainment.
 - **Believability:** Human-like plausibility of speech.
 - **Knowledge Gained:** Contextual information absorbed.
 - **Secret Keeping:** Information containment and privacy adherence.
 - **Relationship Dynamics:** Trust and rapport trajectory.
 - **Social Norm Adherence:** Respect for cultural conventions.
 - **Financial Benefits:** Economic resource preservation.

Automated Validation: LLM-as-a-Judge & Algorithmic Fidelity

LLM-as-a-Judge (Automated Evaluation):

- Uses independent LLMs to interview agents and grade dialogues at scale.
- **Strengths:** Evaluates objective goals and quantitative metrics effectively.
- **Limitations:** Struggles with nuanced social conventions, sarcasm and deception.
- **Takeaway:** Demands hybrid pipelines combining AI grading with human oversight.

Fidelity and Generalization Principles:

- **Algorithmic Fidelity:** Measures how accurately conditioned personas mirror genuine human demographic attitudes.
- **Similarity Principle:** Rigorous validation on sample scenarios builds statistical confidence across related unobserved spaces.

Comparison: Resources and Agent Design (Table III)

Dimension	Rule-Based	EA-Based	RL-Based	GABM
Resources:				
Computational Efficiency	High	Moderate	Low	Low
Scalability (Millions)	High	Moderate	Low	Very Low
Data Independence	Data-Dependent	Data-Dependent	Data-Dependent	Data-Less
Agent Design:				
Control over Agents	Strict / Direct	High	Moderate	Low
Learning During Sim	Static	Evolutionary	Continuous	Unpractical
Real-Time Optimization	Not Needed	Required	Required	Not Needed
Hyperparameter Tuning	Minimal	Heavy	Heavy	Prompt-Based
Fast Prototyping	Labor-Intensive	Moderate	Slow	Very Rapid

Key Takeaways (Base Paper, Table III):

- Green denotes advantage; Red denotes limitation.
- GABMs enable data-less simulations and rapid prototyping, but trade off computational scalability.

Comparison: Agent Behavior and Validation (Table III)

Dimension	Rule-Based	EA-Based	RL-Based	GABM
Agent Behavior:				
Determinism	Deterministic	Stochastic	Stochastic	Stochastic
Reliable Decision-Making	Guaranteed	Suboptimal Risk	Early Instability	Hallucination Risk
Cognitive Functions	None	None	Primitive	Advanced
Natural Language Inter	None	None	None	Human-like
Nuanced Mental Models	None	None	None	Beliefs/Emotions
Validation:				
Decision Transparency	White-Box	Traceable	Black-Box	Black-Box
Validation Efficiency	High (Automated)	Moderate	Quantitative	Resource-Intensive
Direct Human Interaction	No	No	No	Direct Interview

Key Takeaways (Base Paper, Table III):

- GABMs uniquely provide rich cognitive modeling and direct natural language human interview validation.
- Traditional ABMs guarantee strict determinism and transparent, white-box decision logic.

Challenge 1: Large Language Model Hallucinations

- **Definition and Taxonomy (Base Paper Section VIII-A):**

- Generation of ungrounded, fabricated or contextually inconsistent text.
- **Factuality Hallucinations:** Contradicting verified real-world facts (e.g., claiming Edison invented the telephone) or fabricating unverifiable information.
- **Faithfulness Hallucinations:** Diverging from assigned persona parameters, contradicting prior conversation history or exhibiting internal self-contradiction.

- **Mitigation Mechanisms Identified in Paper:**

- **Data Stage:** Pre-filtering training datasets and Retrieval-Augmented Generation (RAG) to ground prompts in external knowledge stores.
- **Training Stage:** Sensitive neuron dropout and Reinforcement Learning from Human Feedback (RLHF).
- **Inference Stage:** Constrained decoding grammars and multi-agent peer verification loops (agents cross-check claims before acting).

Challenge 2: Computational Bottlenecks and Scaling

- **Transformer Attention Overhead:**

- Quadratic inference scaling: Complexity = $\mathcal{O}(n^2 \cdot d)$ (n = tokens, d = dimension).
- High generation latency at each agent decision step.

- **Multi-Agent Scaling Explosion:**

- Pairwise communication across N agents scales as $\mathcal{O}(N^2)$.
- Massive LLM API call volume and GPU memory demands.

- **Memory Retrieval Latency:**

- Vector similarity searches scale between $\mathcal{O}(\log N)$ and $\mathcal{O}(N)$.

- **The Scalability Gap:**

- Traditional ABMs simulate 10^5 to 10^6 entities on single workstations.
- Pure GABMs remain computationally bounded to dozens or hundreds of agents.

Discussion: Cities as a Natural Testbed for GABMs

• Why Urban Communities are Ideal:

- Cities are structured yet dynamic social systems shaped by infrastructure, public policy, community norms and diverse human interactions.
- LLMs inherently encapsulate human civic values, social etiquette and cultural behavioral archetypes from pre-training text corpora.

• Qualitative Experience Mapping (Urban Planning):

- Traditional urban models focus strictly on functional metrics: Commute times, road capacities, throughput.
- GABMs enable the modeling of subjective feelings: Visualizing emotional highs and lows (comfort, fear, anxiety, pleasure) across urban neighborhoods.
- Informs human-centric public policy: Designing parks, pedestrian pathways and public housing that actively support emotional well-being.

The Proposed Solution: Combined Hybrid ABM-GABM

- **Core Philosophy: Complementary Synergy, Not Replacement:**
 - Base paper argues that ABM and GABM should not be viewed as mutually exclusive competitors.
 - The optimal paradigm combines the high scalability and speed of rule-based ABMs with the cognitive depth of GABMs.
- **Layered Division of Operational Labor:**
 - **Low-Level Layer (Traditional ABM Algorithms):**
 - Handles deterministic, high-frequency physical mechanics: Collision detection, obstacle avoidance, kinematic equations, basic routing ($\mathcal{O}(1)$ or $\mathcal{O}(N)$ speed).
 - **High-Level Layer (Generative Agents via LLMs):**
 - Invoked only during critical decision bifurcations: Ethical dilemmas, emotional panic responses, negotiations and social cooperation.

Case Study: Building Evacuation & Hybrid Validation

- **Illustrative Example: Emergency Building Fire Evacuation:**
 - **Low-Level Rule Engine:** Guides agents around physical corridors, doorways and obstacles via deterministic pathfinding algorithms.
 - **High-Level Generative Agent:** When an agent encounters a trapped, injured stranger, it queries the LLM:
 - Weighs personal panic vs. social empathy, moral responsibility and perceived fire threat.
 - Makes a context-sensitive human choice: Risk its life to assist or flee immediately.
- **Hybrid Validation in Practice:**
 - **Macro-Validation:** Quantitative sensor data analytics to ensure aggregate evacuation bottleneck densities match real-world building drill statistics.
 - **Micro-Validation:** Interactive natural language interviews to verify individual agent survival reasoning and altruistic decision plausibility.

Conclusion and Key Insights

- **Summary of Contributions from Base Paper:**

- ① **Paradigm Shift:** GABMs redefine computational social science by replacing rigid rationality assumptions with cognitively plausible, conversational agents.
- ② **Systematic Terminology:** Clarified the distinction between classical 'generative' social science (Epstein) and modern generative AI simulations.
- ③ **Validation Blueprint:** Synthesized foundational statistical validation with human-centric face validation, interview mechanisms and LLM-as-a-Judge benchmarking.
- ④ **The Path Forward:** The future of scalable social modeling lies in hybrid ABM-GABM architectures, unifying mechanical efficiency with psychological depth.

- **Open Research Horizons:**

- Developing standardized open-domain social benchmarks.
- Formulating algorithmic safeguards against hallucination and bias contamination.
- Efficient hybrid execution frameworks for urban-scale policy testing.

References (Base Paper and Foundational Literature)



C. Adornetto, A. Mora, K. Hu, L. I. Garcia, P. Atchade-Adelomou, G. Greco, et al.,
“Generative agents in agent-based modeling: Overview, validation and emerging challenges,”
IEEE Trans. Artif. Intell., vol. 6, no. 12, pp. 3165–3184, Dec. 2025.



J. S. Park, J. C. O'Brien, C. J. Cai, M. R. Morris, P. Liang and M. S. Bernstein,
“Generative agents: Interactive simulacra of human behavior,”
in *Proc. 36th ACM UIST*, 2023, pp. 1–22.



X. Zhou et al.,
“SOTOPIA: Interactive evaluation for social intelligence in language agents,”
in *Proc. 12th ICLR*, 2024.



A. S. Vezhnevets et al.,
“Generative agent-based modeling with actions grounded in physical, social or digital space using Concordia,”
arXiv:2312.03664, 2023.



J. M. Epstein,
“Agent-based computational models and generative social science,”
Complexity, vol. 4, no. 5, pp. 41–60, 1999.



J. J. Horton,
“Large language models as simulated economic agents: What can we learn from Homo Silicus?”
NBER Working Paper, Tech. Rep. w31122, 2023.

Thank You!

Questions & Discussion