

GENERATIVE AGENTS IN AGENT-BASED MODELING: OVERVIEW, VALIDATION, AND EMERGING CHALLENGES

CSQ413 Seminar Presentation

Submitted by:

ARAVIND A KAMATH (CEC23CS041)

Under the Guidance of:

Mrs. JAYASREE K

Assistant Professor

Department of Computer Science and Engineering
College of Engineering, Cherthala

Outline

- 1 Introduction
- 2 ABMs and Generative ABMs
- 3 Evolving Agents
- 4 Generative Agents
- 5 Informal Theory of ABMs Validation
- 6 GABMs Validation
- 7 ABM vs GABM: A Comparison
- 8 GABMs: Emerging Challenges
- 9 Discussion
- 10 Conclusion

Background: Agent-Based Modeling (ABM)

- **Bottom-Up Paradigm:** Simulates autonomous entities (agents) interacting locally within a shared environment.
- **Emergence Principle:** Macro-level patterns emerge naturally from decentralized individual decisions.
- **Wide Applicability:** Used in traffic planning, epidemic spreading and financial market dynamics.
- **Core Objective:** Explains complex collective phenomena without requiring central control.

Limitations of Classical Behavioral Models

- **Strict Economic Rationality:** Assumes agents are perfectly rational utility maximizers (*homo economicus*).
- **Absence of Emotional Nuance:** Cannot capture human stress, fear, empathy or shifting moods.
- **No Natural Language:** Communication is restricted to numbers and binary flags.
- **Rigid Decision Rules:** Fixed if-then heuristics fail when unprecedented situations arise.
- **Population Homogeneity:** Agents are grouped into simple types, ignoring individual diversity.

The Concept of 'Generative' in Social Science

- **Epstein's Principle (1999):** *"If you didn't grow it, you didn't explain it."*
- **Classical Meaning:** "Generative" meant growing macro-structures from simple local rules without central control.
- **Foundational Models:** Schelling's segregation model (1971) and Reynolds' flocking boids (1987).
- **The Technological Shift:** Classical models used hand-coded rules; modern GABMs use Large Language Models for cognitive reasoning.

Traditional ABM vs. Generative ABM: Urban Traffic Flow

Table: Comparison between Traditional ABM and Generative ABM (Table I)

Dimension	Traditional ABM	Generative ABM (GABM)
Scenario	Standard vehicles following signals.	Drivers influenced by stress and habits.
Objective	Maximize throughput and find bottlenecks.	Study how emotions and norms shape traffic.
Outcomes	Local signal timings; hard to generalize.	Generalizable behavioral insights.
Decision Basis	Fixed shortest-path heuristics (Dijkstra).	Contextual reasoning (safety, stress, comfort).

Precursor Agent Architectures: EAs and RL

- **Evolutionary Algorithms (EAs):** Optimize agent rules via selection, crossover and mutation across generations.
- **EA Bottlenecks:** High computational cost, premature local convergence and heavy parameter tuning.
- **Reinforcement Learning (RL):** Optimizes numerical policies to maximize reward signals through trial and error.
- **RL Bottlenecks:** Sample inefficiency, multi-agent instability and total absence of language dialogue.
- **Takeaway:** Neither EAs nor RL provide human-like commonsense reasoning or open-ended communication.

AI Integration in ABM Development




Figure: AI integration across Structural Specification and Output Analysis (Source: Base Paper Figure 1).

Roles of Generative AI in ABM Development

- **Population Synthesis:** Generates diverse agent profiles with realistic ages, traits and occupations.
- **Interactive Decision Making:** LLMs act as cognitive engines, deciding actions dynamically via dialogue.
- **Scenario Generation:** Synthesizes dynamic environments, policy interventions and social events.
- **Macro-Level Tracking:** Uses machine learning to detect patterns and emergent trends across runs.
- **Direct Agent Interviewing:** Allows observers to query agents directly in natural language about their choices.

Defining Generative Agents (GAs)

- **Formal Definition:** Autonomous computational entities powered by Large Language Models capable of human-like behavior and dialogue.
- **Mechanical Layer:** Rule-based physical substrate handling spatial coordinates, collisions and clock time.
- **Cognitive Layer:** LLM processes sensory perceptions, queries memory streams and decides actions.
- **Emergent Behaviors:** Decisions arise from pre-trained world knowledge rather than hand-coded if-then rules.

Internal State: Static Parameters vs. Dynamic Status

- **Static Parameters (Identity Anchor):** Demographic data (age, gender, job), Big Five OCEAN personality traits and social network ties.
- **Dynamic Status (Runtime Representation):** Physical coordinates, current tasks, active goals and emotional affect.
- **Decoupled Architecture:** Separates internal state database storage from the LLM prompt engine.
- **Affective Tracking:** Records sentiment shifts, relationship changes and recent interaction history.

GA Decision Architectures: Feedback Loop and Status




Figure: LLM-based feedback loop architectures for Generative Agents (Source: Base Paper Figure 2).

Inferential vs. Conversational Feedback Loops

- **Inferential Loop [Figure 2(a)]:** Agent receives an environmental event, checks internal status, and the LLM outputs a discrete physical action.
- **Environment Update:** Discrete actions directly alter physical variables in the simulation world.
- **Conversational Loop [Figure 2(b)]:** Agent receives dialogue from peers, updates internal beliefs and status, and outputs natural language text.
- **Social Communication:** Natural language responses directly influence the mental states of neighboring agents.

Core Cognitive Modules: Memory Stream & Retrieval

- **Memory Stream:** Chronological database logging perceptions, dialogues and thoughts in natural language text.

$$\text{Score}(m, q) = \alpha \cdot \text{Recency}(m) + \beta \cdot \text{Importance}(m) + \gamma \cdot \text{Relevance}(m, q)$$

Symbol Definitions:

- $\text{Score}(m, q)$: Total retrieval priority for memory record m under situational query q .
- α, β, γ : Non-negative weights balancing recency, importance and relevance.
- $\text{Recency}(m) = e^{-\lambda \Delta t}$: Exponential decay with decay rate λ and elapsed time Δt .
- $\text{Importance}(m)$: Integer score (1 to 10) assigned by the LLM for event significance.
- $\text{Relevance}(m, q) = \frac{\mathbf{e}_m \cdot \mathbf{e}_q}{\|\mathbf{e}_m\| \|\mathbf{e}_q\|}$: Cosine similarity between embedding vectors $\mathbf{e}_m$ and $\mathbf{e}_q$.

Core Cognitive Modules: Reflection & Dynamic Planning

- **Reflection Hierarchy:** Periodically synthesizes low-level memory records into higher-order abstract thoughts.
- **Belief Tree:** Forms an evolving directed tree from raw observations (leaves) to generalized beliefs (roots).
- **Recursive Planning:** Decomposes broad daily agendas into hourly schedules, then minute-level action sequences.
- **Dynamic Reactive Re-planning:** Pauses routine execution and updates plans when unexpected events occur.

Literature Survey: Existing GABMs (Part 1)

Table: Literature survey of foundational Generative Agent platforms (Source: Base Paper Table II)

Paper & Author	Scenario	Core Architecture	Key Highlights	Validation Approach
Generative Agents <i>Park et al. (2023)</i>	25 agents in 2D RPG Smallville town.	Memory stream + Reflection tree + Daily planning.	Emergent party planning and information spread.	100 human judges across 5 interview categories.
Concordia <i>Vezhnevets et al. (2023)</i>	Grounded physical and social spaces.	Modular language agents with centralized Game Master.	Enforces physical reality and resource constraints.	Qualitative analysis and design guidelines.
SOTOPIA <i>Zhou et al. (2024)</i>	Goal-driven cooperation and negotiation.	Interactive environment with distinct agent profiles.	Quantitative social intelligence benchmarking.	Dual evaluation: Humans vs. LLM-as-a-judge (7 metrics).

Literature Survey: Existing GABMs (Part 2)

Table: Literature survey of social simulation and domain-specific GABMs (Source: Base Paper Table II)

Paper & Author	Scenario	Core Architecture	Key Highlights	Validation Approach
CRSEC Norms <i>CRSEC (2024)</i>	Norm emergence with 10 GAs in Smallville.	4 stages: Creation, Spread, Evaluate, Comply.	Self-governing rules (indoor smoking ban).	30 human evaluators testing compliance rates.
S3 Platform <i>Gao et al. (2023)</i>	Digital social media and opinion dynamics.	LLM agents with attitude and emotion tracking.	Replicates retweet cascades and echo chambers.	Real-world social network dataset matching.
Epidemic GAs <i>Williams et al. (2023)</i>	Disease outbreak and emergency decisions.	Inferential GAs driven by Big Five OCEAN traits.	Personality influence on quarantine compliance.	Empirical matching against epidemic curves.

The Fundamental Validation Dilemma

- **Purpose of Validation:** Ensures simulation models reflect real-world social processes rather than software bugs.
- **White-Box Traditional ABMs:** Model rules and transitions can be inspected and verified mathematically.
- **Black-Box Generative Agents:** Foundation models produce non-deterministic outputs from identical initial prompts.
- **Evidence Hierarchy:** Adapts evidence standards: controlled trials (highest), observational data (middle) and face validation (lowest).

Foundational and Face Validation in Traditional ABM

- **Data Analytics & Docking:** Verifies distributions, correlations and compares outputs against independent models.
- **Empirical Testing & Resampling:** Validates predictions on historical data and uses bootstrapping for confidence bounds.
- **Microface Validation:** Checks individual agent decisions and emotional logic for plausibility.
- **Macroface Validation:** Confirms collective group patterns (crowd flows, segregation) match real-world intuition.
- **Immersive Probing:** Domain experts enter virtual dashboards or VR to observe and interview agents.

GABM Validation Workflows Architecture




Figure: Validation workflows for GABMs across different validation purposes (Source: Base Paper Figure 3).

Integrated Validation Pathways

- **Empirical Pathway (Quantitative Matching):** Conditions initial demographics with census data and tests aggregate trajectories against historical data.
- **Model Alignment:** Applies data analytics and docking across runs to ensure numerical reliability.
- **Macroface Pathway (Collective Patterns):** Validates emergent group dynamics such as spatial clustering and norm spread without central control.
- **Microface Pathway (Individual Realism):** Probes individual agents through structured interviews to check reasoning, memories and consistency.

Automated Validation: LLM-as-a-Judge & Algorithmic Fidelity

- **LLM-as-a-Judge:** Uses independent LLMs to interview agents and grade multi-agent dialogues at large scale.
- **Evaluation Strengths:** High accuracy on objective goals, factual consistency and quantitative criteria.
- **Evaluation Limits:** Struggles with subtle social nuances, passive aggression, sarcasm and deception.
- **Algorithmic Fidelity:** Measures how accurately simulated personas reflect genuine human demographic attitudes.
- **Similarity Principle:** Rigorous testing on sample scenarios builds confidence across related unobserved spaces.

Comparison: Resources and Agent Design

Table: Comparison of ABM Paradigms: Resources and Agent Design (Table III)

Evaluation Dimension	Rule-Based	EA-Based	RL-Based	GABM
Resources				
Computational Efficiency	High	Moderate	Low	Low
Scalability (Millions of Agents)	High	Moderate	Low	Very Low
Data or Domain Knowledge	Dependent	Dependent	Dependent	Data-Less
Agent Design				
Control over Agents	Strict / Direct	High	Moderate	Low
Learning During Simulation	Static	Evolutionary	Continuous	Impractical
Real-Time Optimization	Not Needed	Required	Required	Not Needed
Hyperparameter Tuning	Minimal	Heavy	Heavy	Prompt-Based
Fast Prototyping	Labor-Intensive	Moderate	Slow	Very Rapid

- **Resource Trade-Off:** Rule-based models scale to millions; GABMs require heavy computation.
- **Design Trade-Off:** GABMs enable rapid prototyping with prompts, but offer less direct control.

Comparison: Agent Behavior and Validation

Table: Comparison of ABM Paradigms: Behavior and Validation (Table III)

Evaluation Dimension	Rule-Based	EA-Based	RL-Based	GABM
Agent Behavior				
Determinism	Deterministic	Stochastic	Stochastic	Stochastic
Reliable Decisions	Guaranteed	Suboptimal Risk	Early Instability	Hallucination Risk
Cognitive Functions	None	None	Primitive	Advanced
Natural Language Interaction	None	None	None	Human-like
Nuanced Mental Models	None	None	None	Beliefs / Emotions
Validation				
Decision Transparency	White-Box	Traceable	Black-Box	Black-Box
Validation Efficiency	High	Moderate	Quantitative	Resource-Intensive
Direct Human Interaction	No	No	No	Direct Interview

- **Behavioral Depth:** GABMs uniquely support rich cognitive models and natural language dialogue.
- **Validation Challenge:** Traditional models offer transparent rules; GABMs are opaque black boxes.

Challenge 1: Large Language Model Hallucinations

- **Factuality Hallucinations:** Generating assertions that contradict verified physical environment facts.
- **Faithfulness Hallucinations:** Diverging from assigned persona values, moral codes or prior conversation history.
- **Data Stage Mitigation:** Grounding prompts with Retrieval-Augmented Generation (RAG) using live environment facts.
- **Training Stage Mitigation:** Applying Reinforcement Learning from Human Feedback (RLHF) and persona fine-tuning.
- **Inference Stage Mitigation:** Enforcing constrained grammar decoding (e.g., JSON schema) and multi-agent peer verification.

Challenge 2: Computational Bottlenecks and Scaling

- **Computational Bottlenecks:** Quadratic attention latency and pairwise agent communication scaling:
 $\text{Complexity}_{\text{Attention}} = \mathcal{O}(n^2 \cdot d), \quad \text{Complexity}_{\text{Communication}} = \mathcal{O}(N^2)$

Symbol Definitions:

- n : Sequence length (number of input context tokens in memory prompt).
- d : Embedding representation dimension of the Transformer model.
- N : Total number of interacting generative agents in the simulation.
- M : Number of memory records searched during vector retrieval ($\mathcal{O}(\log M)$ to $\mathcal{O}(M)$).
- **The Scalability Gap:** Traditional ABMs run 10^6 agents; pure GABMs remain limited to dozens or hundreds.

Discussion: Cities as a Natural Testbed for GABMs

- **Alignment with Real-World Text:** Urban laws, historical archives and census data closely match LLM training text.
- **Beyond Physical Metrics:** Traditional urban models measure only functional data (traffic delay, throughput, capacity).
- **Qualitative Experience Mapping (QEM):** Models subjective human feelings: safety perceptions in unlit streets, comfort and transit stress.
- **Human-Centered Planning:** Informs urban policy to support community emotional well-being and pedestrian safety.

The Proposed Solution: Combined Hybrid ABM-GABM

- **Complementary Synergy:** Unifies the high computational speed of rule-based ABMs with the cognitive depth of GABMs.
- **Low-Level Mechanistic Layer (Traditional ABM):** Handles deterministic, high-frequency physics: collision avoidance and kinematics ($\mathcal{O}(1)$ or $\mathcal{O}(N)$).
- **High-Level Cognitive Layer (Generative Agents):** Invoked only during critical decision points: moral choices, negotiations and social cooperation.
- **Key Benefit:** Resolves computational bottlenecks while preserving rich human-like behavior.

Case Study: Building Evacuation & Hybrid Validation

- **Emergency Building Fire Evacuation:** Illustrates the division of labor in an emergency crisis.
- **Low-Level Rule Engine:** Moves agents through corridors, stairwells and exit doors using fast physics heuristics.
- **High-Level LLM Engine:** Decides whether an agent stops to help an injured peer or flees to save itself.
- **Macro-Validation:** Sensor data confirms crowd evacuation clearance times match building drill records.
- **Micro-Validation:** Post-simulation interviews verify individual survival reasoning and moral choices.

Summary of Key Insights and Contributions

- **Cognitive Paradigm Shift:** Replaced rigid rational heuristics with believable, conversational agents.
- **Conceptual Clarity:** Distinguished Epstein's classic generative social science from modern generative AI.
- **Integrated Validation Blueprint:** Formulated a multi-tiered pipeline combining empirical data, face validation and LLM judges.
- **Hybrid Path Forward:** Coupled traditional physics with cognitive LLMs to achieve scalability and psychological realism.

Future Research Trajectories

- **Small Language Models (SLMs):** Fine-tuning compact models to reduce inference latency and computational costs.
- **Standardized Benchmarks:** Establishing open-source multi-agent test environments for objective comparison.
- **Hallucination Bounds:** Formulating mathematical bounds on hallucination rates for policy simulations.
- **Ethical Safeguards:** Preventing bias propagation and harmful stereotypes in simulated human societies.

References



C. Adornetto et al., “Generative agents in agent-based modeling: Overview, validation, and emerging challenges,” *IEEE Trans. Artif. Intell.*, vol. 6, no. 12, pp. 3165-3184, 2025.



J. M. Epstein, “Agent-based computational models and generative social science,” *Complexity*, vol. 4, no. 5, pp. 41-60, 1999.



J. S. Park et al., “Generative agents: Interactive simulacra of human behavior,” in *Proc. 36th ACM UIST*, 2023, pp. 1-22.



A. S. Vezhnevets et al., “Generative agent-based modeling with actions grounded in physical, social or digital space using Concordia,” *arXiv:2312.03664*, 2023.



X. Zhou et al., “SOTOPIA: Interactive evaluation for social intelligence in language agents,” in *Proc. 12th ICLR*, 2024.



C. Gao et al., “S3: Social-network simulation system with large language model-empowered agents,” *arXiv:2307.14984*, 2023.

Thank You!

Questions and Discussion